



# Seq2seq Translation model


CSC401 / 2511 tutorial  
Feb 26, 2020  
Zining Zhu



# Agenda


- Recap: Recurrent models: RNN, LSTM, GRU
- Seq2seq model
  - Setup
  - Train
  - Run
- Beam Search example
- Misc. reminders for A2
- Questions

# RNN




vs.

# LSTM



vs.

# GRU





# RNN models in PyTorch

- For A2, consider RNN, LSTM, and GRU
- The forward() method:
  - Input: hidden\_state, x
  - For the first step, hidden\_state contains all 0 by default
  - hidden\_state for LSTM vs. RNN / GRU
  - x is a tensor of shape (seq\_len, batch\_size, dim), unless you want batch\_first=True
  - Training and eval difference. module.train(), module.eval()
- Check out pytorch docs for details.
  - <https://pytorch.org/docs/stable/nn.html>

# Seq2seq model: Setup

- Encoder + Decoder
- Both nn.Modules




Image source:

<https://towardsdatascience.com/sequence-to-sequence-model-introduction-and-concepts-44121411421>



# Seq2seq model: Encoder

- Encoder
  - Consists of an `nn.Embedding` and a bidirectional RNN (one of LSTM, GRU, and RNN)
  - Input:  $F$  (sequence len  $S$ , batch size  $N$ )
  - The `nn.Embedding` maps input  $F$  onto the rnn input,  $x$  ( $S, N$ , RNN dimension  $I$ )
  - Encoder outputs all hidden states:  $h$  ( $S, N$ , hidden size  $H*2$ )
  - Same for both training and running.



# Seq2seq model: Setup

- Decoder
  - Consists of an RNN / LSTM / GRU and a feed-forward Linear layer.
  - Input / output dependent on training and running status.
- Two versions: without and with attention.
  - DecoderWithoutAttention: Encoder  $\rightarrow h \rightarrow$  first\_hidden\_state
  - DecoderWithAttention: Attend to the  $h$  (encoder outputs).
  - Nevertheless, each decoder step computes log prob of current step.




## Seq2seq model: Train

- “Teacher forcing”
- For each step: given the input and a first hidden state, should learn the correct output (i.e., next token).
  - The input comes from the current token.
  - See `forward()` in `a2_abcs.py` for details.
- Maximum likelihood training
- Train the EncoderDecoder `nn.Module` end-to-end!
  - Gradient passed through both decoder and encoder
  - Remember to `optim.zero_grad()` after `loss.backward()`
  - If computing total loss, use `loss.item()` not `loss` (which preserves the massive graph)



## Sidenote: computation graph in PyTorch

- At each operation, a graph node is created
  - Each tensor contains reference to the graph
- Gradients are computed in `.backward()`
- Some intermediate nodes are discarded
  - BackProp 2nd time error
  - `optim.zero_grad()`
  - See [this notebook](#) for examples.
- The graphs could use up your memory.
  - Graph is gc'ed with tensors.
  - `tensor.detach()` or `tensor.item()`





## Seq2seq model: Run

- “Beam Search Decoding”
- Maintain a “beam” of K partially decoded sentences
- At each step:
  - Compute the likelihood of candidates.
  - Preserve the K candidates with highest likelihood, discard the rest.
- When to stop decoding?
  - When your [top candidate](#) comes to an <eos>
- Go over example in [slides](#) P55-60



## Misc.

- Make sure the shape / types of tensors are correct.
  - Especially in broadcasting / slicing.
  - CPU / GPU tensor types mismatch.
- Debug on small datasets
  - E.g., setting `--max-vocab 100` in building vocabulary. See Appendix A.4 on handout.
- Use a GPU to train the large model.
  - BA GPU labs / AWS EC2 / Google Cloud / etc.,
  - Make sure the code will run on teach.cs -- otherwise will receive 0 marks.
- Ask **and answer others'** questions on piazza
  - But don't post your solution codes.
- Start early! Due March 9 @ 7pm.



## Labs with GPU

- See teaching lab availabilities at <https://www.teach.cs.toronto.edu/faq.html#ABOUT4>
- Sometimes labs are occupied for classes: (dot means “booked”)

time	BA2200	BA3175	BA3185	BA3195
Monday				
10-11am	●	●	●	●
12-2pm				●
2-3pm	●	●	●	●
8-9pm	●	●	●	●
Tuesday				
1-3pm	●	●	●	●
6-8pm	●			
8-9pm		●	●	
Wednesday				
12-2pm		●		
2-4pm	●	●		
4-5pm	●			
6-8pm	●	●	●	●
Thursday				
9-9pm	●	●	●	●
Friday				
10-12pm	●	●	●	●
1-4pm	●	●	●	●
4-6pm	●			